

CSCD 327: Relational Database Systems

Relational Model

Instructor: Dr. Dan Li

Relational Databases

- Most common data model in modern
- Many commercial systems
 - Oracle, MS SQL Server, IBM DB2, more...
- Also open source
 - MySQL, PostgreSQL, more...

Basic Structure

- * Formally, given sets A_1, A_2, \dots, A_n a **relation** r is a subset of

$$A_1 \times A_2 \times \dots \times A_n$$

Thus, a relation is a set of n -tuples (a_1, a_2, \dots, a_n) where each $a_i \in A_i$

- * Example: If

- * $customer_name = \{\text{Jones, Smith, Curry, Lindsay, ...}\}$

- /* Set of all customer names */

- * $customer_street = \{\text{Main, North, Park, ...}\}$ /* set of all street names*/

- * $customer_city = \{\text{Harrison, Rye, Pittsfield, ...}\}$ /* set of all city names */

Then $r = \{$

- (Jones, Main, Harrison),
 - (Smith, North, Rye),
 - (Curry, North, Rye),
 - (Lindsay, Park, Pittsfield) }

is a relation over

$customer_name \times customer_street \times customer_city$

Attribute Types

- Each attribute of a relation has a name
- The set of allowed values for each attribute is called the **domain** of the attribute
- Attribute values are (normally) required to be **atomic**; that is, indivisible
 - E.g. the value of an attribute can be an account number,
but cannot be a set of account numbers
- The special value *null* is a member of every domain
- The null value causes complications in the definition of many operations
 - We shall ignore the effect of null values in our main presentation and consider their effect later

Relation Schema

- A_1, A_2, \dots, A_n are *attributes*
- $R = (A_1, A_2, \dots, A_n)$ is a *relation schema*

Example:

*Customer_schema = (customer_name,
customer_street, customer_city)*

- $r(R)$ denotes a *relation* r on the *relation schema* R

Example:

customer (Customer_schema)

Relation Instance

- * The current values (*relation instance*) of a relation are specified by a table
- * An element t of r is a *tuple*, represented by a *row* in a table
- * Order of tuples is irrelevant (tuples may be stored in an arbitrary order)
- * The degree of a relation is the number of attributes it contains.
- * The cardinality of a relation is the number of tuples it contains.

The diagram shows a table representing a relation instance. The table has three columns and four rows. The columns are labeled *customer_name*, *customer_street*, and *customer_city*. The rows contain the following data: (Jones, Main, Harrison), (Smith, North, Rye), (Curry, North, Rye), and (Lindsay, Park, Pittsfield). Annotations include arrows pointing from the text 'attributes (or columns)' to each of the three column headers, and arrows pointing from the text 'tuples (or rows)' to each of the four rows. The word 'customer' is centered below the table.

<i>customer_name</i>	<i>customer_street</i>	<i>customer_city</i>
Jones	Main	Harrison
Smith	North	Rye
Curry	North	Rye
Lindsay	Park	Pittsfield

customer

Database

- A database consists of multiple relations
- Information about an enterprise is broken up into parts, with each relation storing one part of the information

account : stores information about accounts

depositor : stores information about which customer owns which account

customer : stores information about customers

- Storing all information as a single relation such as
bank(account_number, balance, customer_name, ..)
results in

- repetition of information

- e.g., if two customers own an account (What gets repeated?)

- the need for null values

- e.g., to represent a customer without an account

- Normalization theory deals with how to design relational schemas

Keys

- * Let $K \subseteq R$
- * K is a **superkey** of R if values for K are sufficient to identify a unique tuple of each possible relation $r(R)$
 - * by “possible r ” we mean a relation r that could exist in the enterprise we are modeling.
 - * Example: $\{customer_name, customer_street\}$ and $\{customer_name\}$
are both superkeys of *Customer*, if no two customers can possibly have the same name
 - * In real life, an attribute such as *customer_id* would be used instead of *customer_name* to uniquely identify customers, but we omit it to keep our examples small, and instead assume customer names are unique.

Keys (Cont.)

- * K is a **candidate key** if K is minimal
Example: $\{customer_name\}$ is a candidate key for *Customer*, since it is a superkey and no subset of it is a superkey.
- * **Primary key**: a candidate key chosen as the principal means of identifying tuples within a relation
 - * Should choose an attribute whose value never, or very rarely, changes.
 - * E.g. email address is unique, but may change
- * A relation schema may have an attribute that corresponds to the primary key of another relation. Such attribute is called a **foreign key**.
 - * E.g. *customer_name* and *account_number* attributes of *depositor* are foreign keys to *customer* and *account* respectively.
 - * Only values occurring in the primary key attribute of the **referenced relation** may occur in the foreign key attribute of the **referencing relation**.

SQL Data Definition Language (DDL)

- SQL is primarily a query language, for getting information from a database.
- But SQL also includes a *data-definition* component for describing database schemas.
- **CREATE TABLE** - creates a new database table
- **ALTER TABLE** - alters (changes) a database table
- **DROP TABLE** - deletes a database table

Creating (Declaring) a Relation

- Simplest form is:

```
CREATE TABLE <name> (  
    <list of elements>  
);
```

- To delete a relation:

```
DROP TABLE <name>;
```

Elements of Table Declarations

- Most basic element: an attribute and its type.
- The most common types are:
 - INT or INTEGER (synonyms) - integer
 - REAL or FLOAT- floating point numbers
 - CHAR(n) = fixed-length string of n characters.
 - VARCHAR(n) = variable-length string of up to n characters.

DDL - Primitive Types

- numeric
 - INTEGER (or INT), SMALLINT are subsets of the integers (machine dependent)
 - REAL, DOUBLE PRECISION are floating-point and double-precision floating-point (machine dependent)
 - FLOAT(N) is floating-point with at least N digits
 - DECIMAL(P,D) (or DEC(P,D), or NUMERIC(P,D)):
 - P: The maximum total number of decimal digits that can be stored, both to the left and to the right of the decimal point.
 - D: The maximum number of decimal digits that can be stored to the right of the decimal point.

DDL - Primitive Types (cont.)

- character-string
 - CHAR(N) (or CHARACTER(N)) is a fixed-length character string
 - VARCHAR(N) (or CHAR VARYING(N), or CHARACTER VARYING(N)) is a variable-length character string with at most N characters
- bit-strings
 - BIT(N) is a fixed-length bit string
 - VARBIT(N) (or BIT VARYING(N)) is a bit string with at most N bits

DDL - Primitive Types (cont.)

- DATE and TIME are types in SQL.

- The form of a date value is:

yyyy-mm-dd

- Example: DATE '2007-09-30' for Sept. 30, 2007.

- The form of a time value is:

hh:mm:ss

with an optional decimal point and fractions of a second following.

- Example: TIME '15:30:02.5' = two and a half seconds after 3:30PM.

Example: Create Table

```
CREATE TABLE account
    (account_int      varchar(15) ,
     branch_name
    varchar(15) ,
     balance          int
    ) ;
```


Declaring Keys

- An attribute or list of attributes may be declared PRIMARY KEY or UNIQUE (more about unique from chapter 11).
- Either says that no two tuples of the relation may agree in all the attribute(s) on the list.
- Place PRIMARY KEY or UNIQUE after the type in the declaration of the attribute.
- Example:

```
CREATE TABLE account
  (account_int      varchar(15) unique,
   branch_name      varchar(15),
   balance          int,
   primary key (account_int) );
```

PRIMARY KEY vs. UNIQUE

1. There can be only one PRIMARY KEY for a relation, but several UNIQUE attributes.
2. No attribute of a PRIMARY KEY can ever be NULL in any tuple. But attributes declared UNIQUE may have NULLs, and there may be several tuples with NULL.

Changing a Relation

- Simplest form is:
 ALTER TABLE <name>
 action;
- The SQL standard restricts each ALTER TABLE statement to a **single** table change.
- Add a column definition to a table
- Drop a column from a table
- Change the default value for a column
- Add or drop a primary key for a table
- Add or drop a new foreign key for a table
- Add or drop a uniqueness constraint for a table
- Add or drop a check constraint for a table

Example: Alter Table

```
ALTER TABLE CUSTOMER
```

```
    ADD CONTACT_PHONE CHAR(10);
```

The new columns will have NULL values for existing customers.

```
ALTER TABLE CUSTOME
```

```
    ADD PRIMARY KEY (customer_name);
```