

# CSCD 327: Relational Database Systems

## Multitable Queries

Instructor: Dr. Dan Li

# An Example

- *List all orders, showing the order number and amount, and the name and credit limit of the customer who placed it.*
  - Order number and amount come from ORDERS
  - Customer name and credit limit come from CUSTOMERS
  - There is a connection between ORDERS and CUSTOMERS
    - CUST\_NUM in CUSTOMERS is a primary key
    - CUST in ORDERS is a foreign key which refers to CUST\_NUM in CUSTOMERS
    - They build a PARENT (the table containing the primary key) / CHILD (the table containing the foreign key) relationship
- How would you get the query results by hand?

# Two Findings

- Each row of query results draws its data from a specific *pair of rows, one from the ORDERS table and one from the CUSTOMERS table.*
  - Need to select from TWO tables
- The pairs of rows are found by matching the data values in *corresponding columns* from the tables.
  - Need to specify such condition in WHERE clause
- Solution

```
SELECT ORDER_NUM, AMOUNT, COMPANY, CREDIT_LIMIT
FROM ORDERS, CUSTOMERS
WHERE CUST = CUST_NUM;
```

# Parent/Child Relationship

- Parent: the table containing the primary key (CUSTOMER)
- Child: the table containing the foreign key (ORDER)
- One-to-many relationship
  - Each order is associated with exactly one customer.
  - Each customer can have many associated orders.
- Identify parent and child
  - List each salesperson and the city and region where they work.

```
SELECT NAME, CITY, REGION  
FROM SALESPERSONS, OFFICES  
WHERE REP_OFFICE = OFFICE;
```

SALESPERSONS is the child;  
OFFICES is the parent.

- List the offices and the names and titles of their managers.

```
SELECT CITY, NAME, TITLE  
FROM OFFICES, SALESPERSONS  
WHERE MGR = EMPL_NUM;
```

SALESPERSONS is the parent;  
OFFICES is the child.

# JOIN ... ON to Replace WHERE

- List each salesperson and the city and region where they work.

```
SELECT NAME, CITY, REGION  
FROM SALESREPS JOIN OFFICES  
ON REP_OFFICE = OFFICE;
```

- List the offices and the names and titles of their managers.

```
SELECT CITY, NAME, TITLE  
FROM OFFICES JOIN SALESREPS  
ON MGR = EMPL_NUM;
```

# Joins with Row Selection Criteria

- *List the offices with a target over \$600,000 and their manager information.*

```
SELECT CITY, NAME, TITLE  
FROM OFFICES, SALESREPS  
WHERE MGR = EMPL_NUM  
AND TARGET > 600000.00;
```

Same as:

```
SELECT CITY, NAME, TITLE  
FROM OFFICES JOIN SALESREPS  
ON MGR = EMPL_NUM  
WHERE TARGET > 600000.00;
```

# Multiple Parent/Child Relationships

- *List all the orders, showing amounts and product descriptions.*

```
SELECT ORDER_NUM, AMOUNT, DESCRIPTION  
FROM ORDERS, PRODUCTS  
WHERE MFR = MFR_ID  
AND PRODUCT = PRODUCT_ID;
```

Same as:

```
SELECT ORDER_NUM, AMOUNT, DESCRIPTION  
FROM ORDERS JOIN PRODUCTS  
ON MFR = MFR_ID  
AND PRODUCT = PRODUCT_ID;
```

# NATURAL JOIN

- Can only be used when the two matching columns have the exactly same name in both tables.
  - Not the case in previous examples
  - Now let's make some changes
    - Change MFR\_ID in PRODUCTS into MFR
    - Change Product\_ID in PRODUCTS into Product
- Match ALL column pairs with the same names.
- All of the following four statements are the same.

```
SELECT ORDER_NUM, AMOUNT, DESCRIPTION  
FROM ORDERS NATURAL JOIN PRODUCTS;
```

```
SELECT ORDER_NUM, AMOUNT, DESCRIPTION  
FROM ORDERS JOIN PRODUCTS  
USING (MFR, PRODUCT);
```

```
SELECT ORDER_NUM, AMOUNT, DESCRIPTION  
FROM ORDERS JOIN PRODUCTS  
ON ORDERS.MFR = PRODUCTS.MFR  
AND ORDERS.PRODUCT = PRODUCTS.PRODUCT;
```

```
SELECT ORDER_NUM, AMOUNT, DESCRIPTION  
FROM ORDERS, PRODUCTS  
WHERE ORDERS.MFR = PRODUCTS.MFR  
AND ORDERS.PRODUCT = PRODUCTS.PRODUCT;
```

Q: which one is preferred?

A: JOIN... USING

Why?



# JOIN Involving Three or More Tables

- *List orders over \$25,000, including the name of the salesperson who took the order and the company name of the customer who placed it.*

```
SELECT ORDER_NUM, AMOUNT, COMPANY, NAME
FROM ORDERS, CUSTOMERS, SALESREPS
WHERE CUST = CUST_NUM
AND REP = EMPL_NUM
AND AMOUNT > 25000.00;
```

CUST in ORDERS is a foreign key to the CUSTOMERS table (CUST\_NUM);  
REP in ORDERS is a foreign key to the SALESREPS table (EMPL\_NUM).

Same as:

```
SELECT ORDER_NUM, AMOUNT, COMPANY, NAME
FROM ORDERS JOIN CUSTOMERS ON CUST = CUST_NUM
JOIN SALESREPS ON REP = EMPL_NUM
WHERE AMOUNT > 25000.00;
```

# Is Parent/Child Relationship a Must?

- NO.
- Any columns can serve as matching columns, provided they have comparable data types.
- *Find all orders received on a day when a new salesperson was hired.*

```
SELECT ORDER_NUM, AMOUNT, ORDER_DATE, NAME  
FROM ORDERS, SALESREPS  
WHERE ORDER_DATE = HIRE_DATE;
```

This example generates a many-to-many relationship:  
Each potential order\_date can be associated with more new hires.  
Each potential hire\_date can be associated with more orders.

# Is Equality Condition a Must?

- No
- *List all combinations of salespeople and offices where the salesperson's quota is more than that office's target, regardless of whether the salesperson works there.*

```
SELECT NAME, QUOTA, CITY, TARGET FROM  
SALESREPS, OFFICES  
WHERE QUOTA > TARGET;
```

# Qualified Column Names to Avoid Ambiguity

- *Show the name, sales, and office for each salesperson.*

```
SELECT NAME, SALES, CITY  
FROM SALESREPS, OFFICES  
WHERE REP_OFFICE = OFFICE;
```

It won't work, because  
SALES appear in both  
tables.

- To avoid ambiguity, use dot (.) operator to explicitly list table name.

```
SELECT NAME, SALESREPS.SALES, CITY  
FROM SALESREPS, OFFICES  
WHERE REP_OFFICE = OFFICE;
```

# Self-Joins

- *List the names of salespeople and their managers.*

```
SELECT NAME, NAME  
FROM SALESREPS, SALESREPS  
WHERE MANAGER = EMPL_NUM;
```

This SELECT statement is illegal because of the duplicate reference to the SALESREPS table in the FROM clause.

How about eliminating one reference?

```
SELECT NAME, NAME  
FROM SALESREPS  
WHERE MANAGER = EMPL_NUM;
```

This query is legal, but it won't do what you want it to do!

Correct solution:

```
SELECT EMPS.NAME, MGRS.NAME  
FROM SALESREPS EMPS, SALESREPS MGRS  
WHERE EMPS.MANAGER = MGRS.EMPL_NUM;
```

Use aliases to refer to the same table twice.  
Can be considered as two virtual copies.  
To simplify, only use one alias; the other one uses the original table name.

# Multitable Query Processing

- The query is processed in this order:
  - Form the product of the tables named in the FROM clause.
  - Apply matching-column condition specified by ON clause.
  - Apply select condition specified by WHERE clause.
  - Keep the columns listed in SELECT clause.
  - Remove duplicates if SELECT DISTINCT is specified.
  - Sort the query results based on ORDER BY clause.

# More JOINS

- OUTER JOIN to avoid loss of information
  - Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
  - Uses *null* values.
  - LEFT OUTER JOIN
  - RIGHT OUTER JOIN
  - FULL OUTER JOIN (MySQL doesn't support it!)
- INNER JOIN

# Joined Relations

- **Join operations** take two relations and return as a result another relation.
- These additional operations are typically used as subquery expressions in the **from** clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

## *Join types*

inner join  
left outer join  
right outer join  
full outer join

## *Join Conditions*

natural  
on <predicate>  
using ( $A_1, A_1, \dots, A_n$ )



# Sample Relations

- Relation *course*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

- Relation *prereq*

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

- Observe that

prereq information is missing for CS-315 and  
course information is missing for CS-347

# Left Outer Join

- *course* **natural left outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>

# Right Outer Join

■ *course* **natural right outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

# Full Outer Join

- *course* **natural full outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

# Inner Join & ON

- *course* **inner join** *prereq* **on**  
*course.course\_id = prereq.course\_id*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

- What is the difference between the above, and a natural join?
- *course* **left outer join** *prereq* **on**  
*course.course\_id = prereq.course\_id*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	<i>null</i>	<i>null</i>

# USING

- *course* **full outer join** *prereq* **using** (*course\_id*)

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101