

# Last Class

- 1, Two ways to define shared memory
- 2, Common Cuda Programming Strategy
- 3, Memory Access Efficiency, CGMA

# Two ways to Define Shared Mem

- Two ways of defining shared memory

## **Dynamic (second kind)**

// when the size of the array isn't known at compile time...

```
__global__ void adj_diff(int *result, int *input)
```

```
{
```

```
    // use extern to indicate a __shared__ array will be
```

```
    // allocated dynamically at kernel launch time
```

```
    extern __shared__ int s_data[]; // This has to be one dimensional. If you use 2D tile, you have to  
                                     // flatten the 2D tile into 1D shared memory.
```

```
    ...
```

```
}
```

// pass the size of the per-block array, in bytes, as the third

// argument to the triple chevrons,

// used with deviceQuery functions.

```
adj_diff<<<num_blocks, block_size, block_size * sizeof(int)>>>(r,i);
```


# Today

- 1, Strategy to reduce global memory traffic
- 2, Tiled Matrix-matrix multiplication kernel
- 3, Analyzing how much global memory access is reduced for the new kernel?

# Small Shared Mem

- Global Memory is big, but very slow.
- Shared memory is 100X faster, but small.
- Strategy to use Shared Memory
  - Partition data into tile
  - Each tile can fit into shared memory.
  - Tiles should be independent to each other when they are being processed.
  - Reduce Global Memory Traffic


# Strategy to Reduce Global Mem Traffic



**FIGURE 5.5**

A small example of matrix multiplication. For brevity, We show  $d\_M[y*Width + x]$ ,  $d\_N[y*Width + x]$ ,  $d\_P[y*Width + x]$  as  $M_{y,x}$ ,  $N_{y,x}$ ,  $P_{y,x}$ , respectively.

# Strategy to Reduce Global Mem Traffic


Access order 

thread <sub>0,0</sub>	$M_{0,0} * N_{0,0}$	$M_{0,1} * N_{1,0}$	$M_{0,2} * N_{2,0}$	$M_{0,3} * N_{3,0}$
thread <sub>0,1</sub>	$M_{0,0} * N_{0,1}$	$M_{0,1} * N_{1,1}$	$M_{0,2} * N_{2,1}$	$M_{0,3} * N_{3,1}$
thread <sub>1,0</sub>	$M_{1,0} * N_{0,0}$	$M_{1,1} * N_{1,0}$	$M_{1,2} * N_{2,0}$	$M_{1,3} * N_{3,0}$
thread <sub>1,1</sub>	$M_{1,0} * N_{0,1}$	$M_{1,1} * N_{1,1}$	$M_{1,2} * N_{2,1}$	$M_{1,3} * N_{3,1}$

FIGURE 5.6

Global memory accesses performed by threads in block<sub>0,0</sub>.

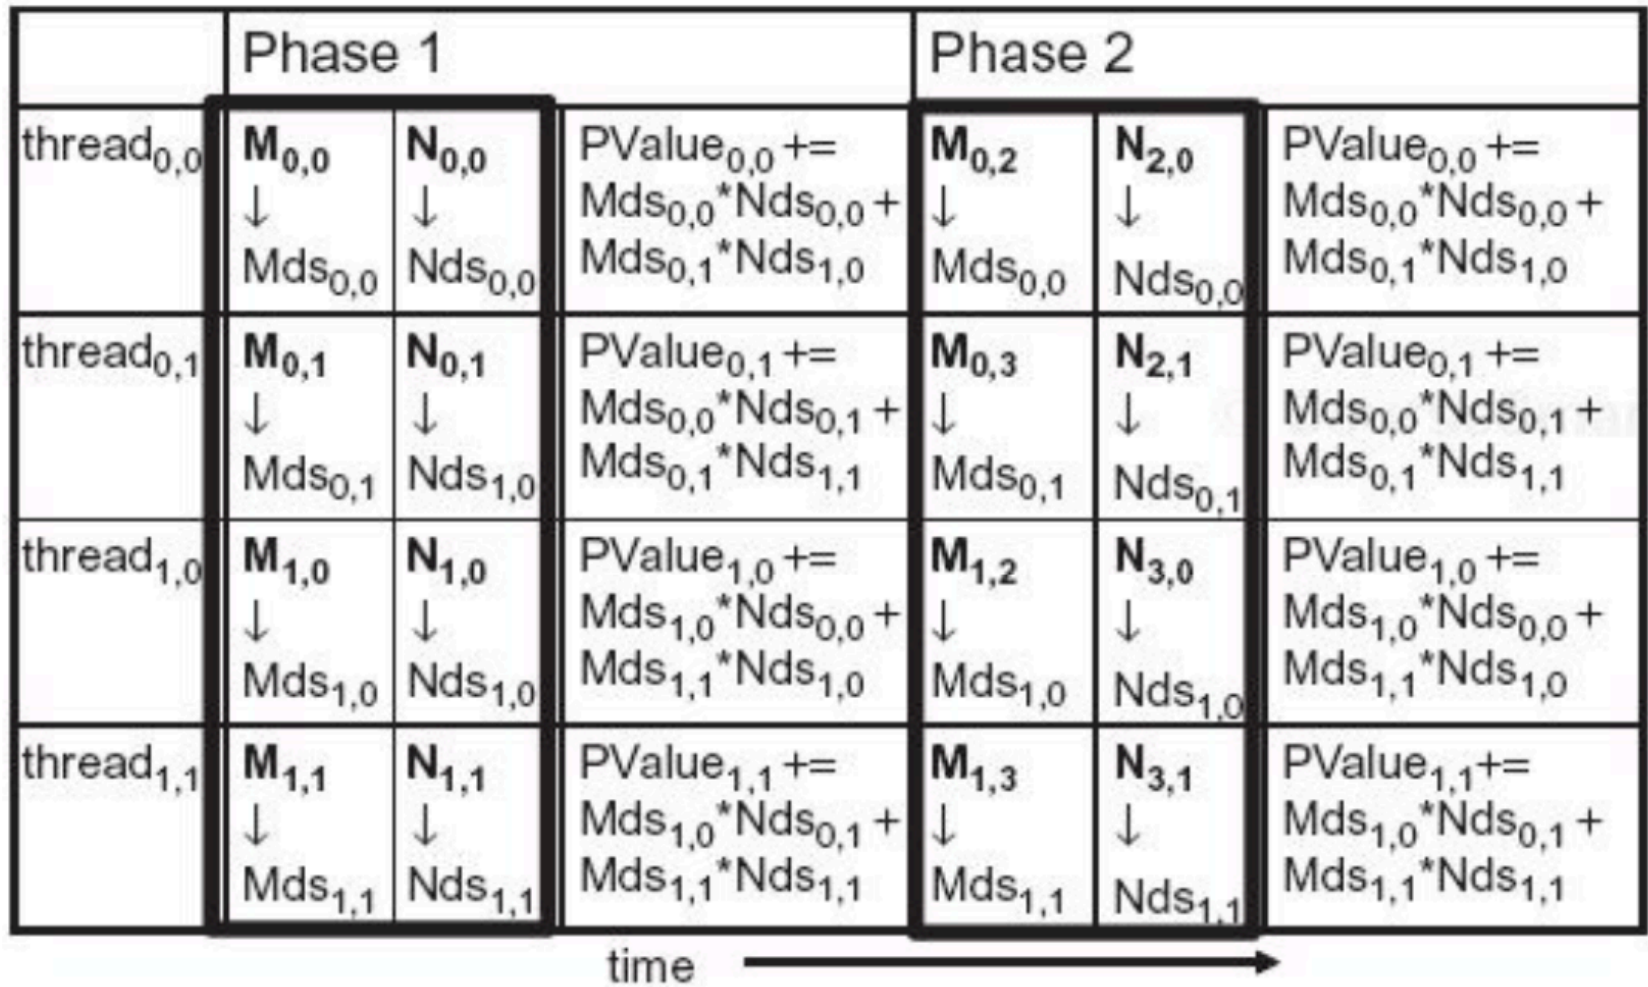
# Tiled Matrix-Matrix Multiplication



**FIGURE 5.10**

Tiling  $M$  and  $N$  matrices to utilize shared memory.

# Tiled Matrix-Matrix Multiplication



**FIGURE 5.11**

Execution phases of a tiled matrix multiplication.



```

__global__ void MatrixMulKernel(float* d_M, float* d_N, float* d_P,
    int Width) {

1.  __shared__ float Mds[TILE_WIDTH][TILE_WIDTH];
2.  __shared__ float Nds[TILE_WIDTH][TILE_WIDTH];

3.  int bx = blockIdx.x;  int by = blockIdx.y;
4.  int tx = threadIdx.x; int ty = threadIdx.y;

    // Identify the row and column of the d_P element to work on
5.  int Row = by * TILE_WIDTH + ty;
6.  int Col = bx * TILE_WIDTH + tx;

7.  float Pvalue = 0;
    // Loop over the d_M and d_N tiles required to compute d_P element
8.  for (int m = 0; m < Width/TILE_WIDTH; ++m) {

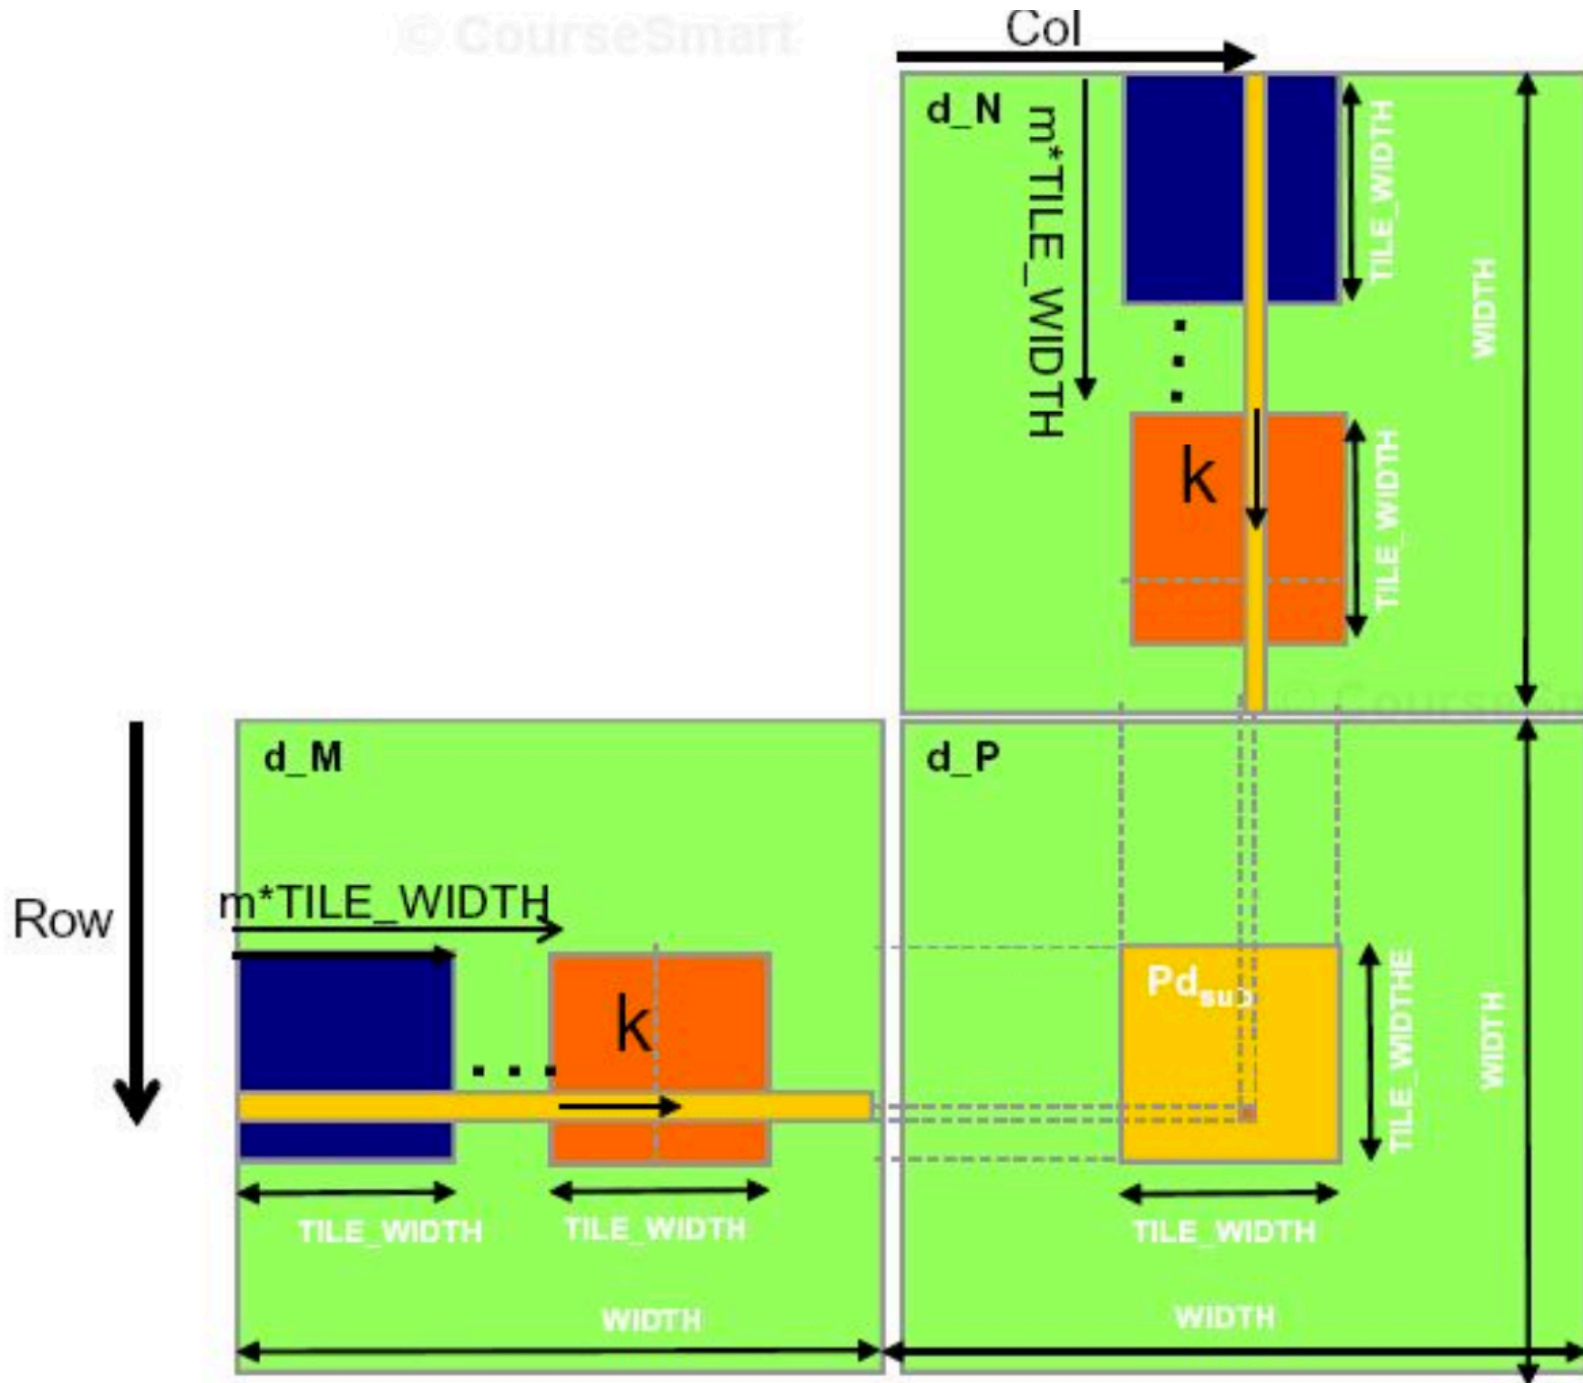
        // Collaborative loading of d_M and d_N tiles into shared memory
9.      Mds[ty][tx] = d_M[Row*Width + m*TILE_WIDTH + tx];
10.     Nds[ty][tx] = d_N[(m*TILE_WIDTH + ty)*Width + Col];
11.     __syncthreads();

12.     for (int k = 0; k < TILE_WIDTH; ++k) {
13.         Pvalue += Mds[ty][k] * Nds[k][tx];
14.     }
15.     __syncthreads();
    }
    d_P[Row*Width + Col] = Pvalue;
}

```

**FIGURE 5.12**

Tiled matrix multiplication kernel using shared memory.



# Analysis

- If we use `TILE_WIDTH=16`, also the block size is  $16 \times 16$ .
- We can reduce the global memory accesses by a factor of 16.
  - Increase the CGMA from 1 to 16.
  - Improve throughput of the Cuda device.

# Analysis

- If we use  $\text{TILE\_WIDTH}=16$ , also the block size is  $16 \times 16$ .
- We can reduce the global memory accesses by a factor of 16. How to calculate?
  - In figure 5.13, focus on the tile  $\text{Pd}_{\text{sub}}$  in  $\text{d\_P}$ ,
  - For simple Matrix-Matrix Multiplication Kernel, total global access  $\text{C1} = (\text{WIDTH} * \text{TILE\_WIDTH}) * \text{TILE\_WIDTH} * 2$
  - For tiled Matrix-Matrix Multiplication Kernel, total global mem access  $\text{C2} = (\text{WIDTH} * \text{TILE\_WIDTH}) * 2$
  - $\text{C1} / \text{C2} = \text{TILE\_WIDTH}$

# Analysis

- If our device has a 200GB/s global memory bandwidth, with CGMA = 16,
- Now it supports  $(200 / 4) * 16 = 800$  GFLOPS.
  - Close to its peak performance.

# Wrap up

- 1, Strategy to reduce global memory traffic
- 2, Tiled Matrix-matrix multiplication kernel
- 3, Analyzing how much global memory access is reduced for the new kernel?

# Next Class

- 1, Memory as a limiting factor to parallelism
- 2, Start chapter 6, Warp and Thread Execution